

exML: Explainable maximum likelihood tool for analysis of DNA samples

ABSTRACT

We describe *exML*, an explainable analysis tool for DNA classification. The underlying problem is to estimate the proportion of different species that contribute to a given genomic dataset. While there are available solutions, they lack explainability. To this end, our proposed solution incorporates novel explanation methods that accompany a Maximum Likelihood analysis in this context. The explanations the tool generates include a variant of the commonly used attribution method that assigns influence scores to individual datums based on Shapley values, and a notion of counterfactuals that identify subsets of the input data that together are most influential. As a use case for this paper, we will focus on the analysis of ancient DNA samples collected from archaeological sites, and show how the explanations generated by *exML* provide insights on otherwise ambiguous classification results.

1 INTRODUCTION

The amount of genomic data has been increasing dramatically and steadily ever since the first full human genome was sequenced in 2001 [6], to a state that it is predicted that by the end of 2025, genomes of 60 million humans will be sequenced and analyzed [1]. This allowed research progress that facilitates advances in the fields of healthcare, evolutionary research, and forensic investigations.

One main algorithmic challenge in the field, which will be the focus of this paper, is to quantify the contribution of different organisms to a given DNA dataset [5]. Formally, given a set of species S and a dataset of DNA reads D , the goal is to output a distribution vector that assigns for each $s \in S$ a proportion that estimates how abundant s is in D . Though there are existing methods to the problem that generally attain good accuracy, they often perform badly when used to classify small datasets (<100 DNA reads), and lack explainability (see [8]) which may lead to errors and difficulties in resolving ambiguities in the model output.

As a use case, we will focus on tackling this classification problem on data of *ancient DNA of hominins*; Sequencing the DNA of ancient individuals can be used as a virtual time machine to explore our evolutionary history and that of our closest extinct relatives, the Neanderthals and the Denisovans. To date, over 6,000 ancient human genomes have been sequenced from fossilized skeletal remains [3], to assemble a database of labeled reference sequences. The retrieval of DNA from sediments deposited at archaeological sites has recently opened new avenues of research (e.g. [7]), allowing to generate ancient genetic data even in the absence of skeletal remains, by extracting and sequencing the DNA molecules that were deposited in the sediment for up to tens of thousands of years. Yet, this new type of data is inherently more complicated to analyze, requiring the development of new tools, to deal with the challenges of deducing research insights from *small amounts of data, corruption of DNA molecules* over tens of thousands of years, and the fact that a single DNA data sample can have DNA sequences *originate from multiple individuals*.

Explainability is a key challenge in this respect as the outputted proportions may, in some cases, be ambiguous. For example, if a system estimates that a certain dataset consists of 50% Neanderthals and 50% Homo sapiens, this result can be interpreted in different ways: 1) Neanderthals and Homo sapiens lived jointly in a single location. 2) All reads are ambiguous, leading to the inability of the model to determine the species of origin. 3) The data originated from another species that is equally distant to both Neanderthals and Homo sapiens. 4) 10% of the data is from Homo Sapiens, 10% from Neanderthals, and the rest is unidentified. State-of-the-art solutions do not allow to easily differentiate between the three, whereas we will exemplify how the system we describe in the paper, and specifically the versatile explanations it generates, can shed light on the correct way to interpret such output.

We present exML, an explainable DNA analysis system that addresses the above challenges. The system is both attaining state-of-the-art level accuracy in the classification problem, and also provides multiple explanations to its output. The system contains a Maximum Likelihood algorithm for the proportion estimation, and is also equipped with explainability capabilities of multiple flavors, that are on the one hand tailored to the specific settings but may be of general interest in the context of explaining Maximum Likelihood systems. We identify two axes of explanations, the first is the type of explanation, namely (a) explanations that are based on quantifying the contribution of individual datums (e.g., DNA reads) or (b) explanations that are based on counterfactuals, which are dataset modification that affect the classification in certain ways. The second axis is whether we explain a classification via data points from the given dataset or via hyper-parameters of the model; the latter are assigned in our setting using a database of *DNA references* which are labeled full genomes of an organism.

In designing each type of explanation, we are inspired by existing methods, but adapt them to our settings. Two kinds of adaptations are in order: First, we need to adapt existing definitions. For instance, in the context of attribution of individual contribution, a standard solution is SHAP [4] which is based on Shapley values. We show that in our context, a modification to Shapley values is needed, as they are based on quantifying the influence of a player on the *result* of a game, when being added to a subset of players, whereas we estimate the influence of adding a DNA read on the *proportions* of species. This requires to take into account that in our case, the influence of adding a single player to a group is expected to be smaller the bigger the group is. Another adaptation we applied is to adjust the explanation computations to our settings. For instance, both Shapley-based and Counterfactual-based explanations involve executing the explained algorithm on multiple subsets of the data, which is typically computationally expensive. In our settings however, we were able to optimize this computation by utilizing the structure of the Maximum Likelihood method.

In this paper, we describe the maximum likelihood algorithm we applied on the classification problem, and show the accuracy it

attains on several datasets consisting of simulated and real-world ancient DNA reads of hominids (Homo sapiens, Neanderthals, or Denisovan). We also demonstrate how the different explanations exML generates complement each other and allow users to obtain a holistic view of the analysis results and gain further insights on the data and model.

2 ALGORITHM FOR DNA CLASSIFICATION

We first describe the general settings, and a maximum likelihood algorithm for the problem. We have in hand a dataset D of DNA reads and a dataset R of labeled DNA references. Each of the references in R is labeled with an item of a group S that intuitively corresponds to biological species. Our goal is to estimate the proportions of the different species in D , formally, to output a distribution vector v , such that for each $s \in S$, v_s is the estimated proportion of species s in D (of course $\sum_i v_i = 1, v_i \geq 0$). We further have a substitution matrix M s.t $M_{i,j,k}$ is the likelihood of observing the letter j in index k of a DNA read, given that the DNA read is originated from an organism that is closely related to an organism that has the letter i in index k . This matrix is used to estimate the likelihood that an observed DNA read d is related to a reference r , and intuitively addresses the facts that genomes of closely related organisms are very similar, but not identical, and the DNA corruption that can affect the input dataset D . Table 1 summarizes the main notations.

Algorithm 1 is a Maximum Likelihood estimator for this task. It consists of two steps (captured by methods invocation in lines 1 and 2): PreProcessing, namely estimating the likelihood of the observed data as a function of the model parameters, and MaximizeLikelihood, namely finding parameters to optimize this function. The PreProcessing method first aligns every DNA read d to every labeled DNA reference r , to find r_d - sub sequence of r that is most similar to d (line 6). Then, using r_d and M , it calculates $p[r_d]$ - the estimated likelihood to observe the DNA read d in a genome of an organism that is closely related to an organism with genome r (line 7). In lines 9-11, for each species $s \in S$ and read $d \in D$, it calculates $A_{s,d}$ - the average of the $P[r_d]$ values for references labeled with s . The PreProcessing method returns the $A_{s,d}$ values, and they are sent as parameters to MaximizeLikelihood, that uses them to return the species proportions that maximizes the likelihood of D (lines 15-16), which is the output of the algorithm. Note that the algorithm is highly parallelizable, as the entire loop in lines 4-11 can be executed independently for every DNA read d . Optimization in line 16 is done using a variant of the gradient descent algorithm - start with a random distribution vector, and in every iteration, make a small step to the direction that maximizes the increase in the likelihood value.

We implemented Algorithm 1 and tested its performance on 4 datasets of ancient human DNA reads collected from sediments, and on simulated datasets of ancient hominin DNA, in which the correct species proportion is known (data is simulated based on real genomes of hominins using the ancient DNA simulator tool published in [7]). For each dataset, we executed Algorithm 1 and calculated the KL divergence [2] between the correct proportions and the output of the algorithm. Figure 1 shows the results and compares our tool to Kallisto, a method based on k-mers and pseudo alignment that is the state-of-the-art method in classification of




Figure 1: Comparison of exML and kallisto method, we generated 500 samples for each sample size, and calculated the average KL divergence between the output of the different algorithms and the correct answer (smaller KL divergence from correct answer means better accuracy).

DNA read	A word $\in \{A, C, T, G\}^*$
D	Input dataset of DNA reads
S	Set of possible labels (species)
R	Set of DNA references (full genomes)
$Y \in S^{ R }$	Labels for references, $Y_i \in S$ is the label of R_i
$\alpha^* \in [0, 1]^{ S }$	Algorithm's output - distribution vector ($\sum_i \alpha_i^* = 1$ and $\alpha_i^* \geq 0$)
R_s	$\{R_i \in R Y_i = s\}$
M	Substitution matrix
$v(D) \in [0, 1]^{ S }$	Result of running Algorithm 1 on dataset D .

Table 1: Table of notations

ancient DNA (see [7]). Note that when $|D|$ is small (as usually is the case in ancient DNA data), this simple algorithm outperforms state-of-the-art technique, and in larger $|D|$'s it is at least as accurate (though it runs slower, due to the computationally expensive alignment step).

Example 2.1. We generated 2 datasets on which we will exemplify the output of Algorithm 1: D_1 contains 20 DNA reads from Homo sapiens, 20 DNA reads from Denisovans, and 40 DNA reads from regions in the genome that are conserved across all hominin genomes (*i.e.*, non informative regions). D_2 contains 40 DNA reads from Homo sapiens and 40 DNA reads from Denisovans. The average length of a DNA read in both datasets is 75. When running Algorithm 1 on these datasets, using the label set $\{Homo sapiens, Neanderthal, Denisovan\}$, and a database of labeled genomes of hominin species as references set, the output was $v(D_1) = v(D_2) = (0.5, 0, 0.5)$. This output intuitively means that for both datasets, the estimation of the system is that the proportions that best fit the data are 50% Homo sapiens DNA, no Neanderthal DNA and 50% Denisovan DNA.

Algorithm 1: Maximum likelihood algorithm to estimate proportion of species in genomic dataset

Data: input dataset D , substitution matrix M , label set S , labeled set of DNA references R

Result: Distribution α^* - species proportion estimation

```

1  $A = \text{PreProcessing}(D, R, M)$ 
2 return  $\text{MaximizeLikelihood}(A, D)$ 
3 Procedure  $\text{PreProcessing}(D, R, M)$ 
4   foreach  $d \in D$  do
5     foreach  $r \in R$  do
6       Align  $d$  to  $r$  to get  $r_d :=$  sub sequence of  $r$  that is
        most similar to  $d$ 
7        $P[r_d] = \prod_{t=1}^{|d|} M(r_d[t], d[t], t)$ 
8     end
9     foreach  $s \in S$  do
10       $A_{s,d} = \frac{\sum_{r \in R_s} P[r_d]}{|R_s|}$ 
11    end
12  end
13  return all  $A_{s,d}$  values
14 Procedure  $\text{MaximizeLikelihood}(A, D)$ 
15   $L(D; \alpha) := \prod_{d \in D} \sum_{s \in S} \alpha_s * A_{s,d}$ 
16  return  $\alpha^* := \text{argmax}_{\alpha} (L(D; \alpha))$  with the constraints
     $\sum_i \alpha_i = 1$  and  $\alpha_i \geq 0 \ \forall i$ .
```

3 EXPLANATIONS

In the explanations our system produces, we employ state-of-the-art techniques for explanations, and adapt them to the particular settings in hand. Specifically, we observe that explanations are relevant for both *reads* and *references*. We next detail both types, describe the algorithms our system utilizes to generate them, and show examples to their usefulness on real world scenarios in analysis of ancient DNA.

3.1 Read-level explanations

A read-level explanation aims to quantify and explain the influence of a single read, or a subset of reads, from the input data set on the output of the algorithm. We present two kinds of explanations at the read-level: attribution-based and Counterfactuals.

3.1.1 Shapley based read attribution.

Definition 3.1. Given a set of DNA reads D , a set of labels S , and an algorithm A that gets D as input and outputs a distribution vector over S , we define a *read-level attribution* for $d \in D$ as a vector $v^d \in R^{|S|}$, such that v_s^d is the influence of d on the proportion A assigns to label $s \in S$.

Example 3.2. As an example for definition 3.1, A can be Algorithm 1 described above, that gets as input a dataset D of DNA reads, S is the set $\{\text{HomoSapiens}, \text{Neanderthal}, \text{Denisovan}\}$ and the explanation of a single DNA read d is a vector $v^d \in R^3$ where v_0 quantifies how d changes the proportion Algorithm 1 assigns to Homo sapiens. For instance, if v_0^d is a positive number, it means that d makes the algorithm tend to increase the estimated proportion of Homo Sapiens.

To generate attribution based read-level explanations, we apply the notion of Shapley values, with modifications to our settings. In the context of cooperative games, Shapley values estimate the contribution of a single player to the result of the game. The Shapley value of a player i is defined as (v is the value of a game and N is the number of players):

$$\frac{1}{N!} \sum_{G \subseteq [N] \setminus \{i\}} [v(G \cup i) - v(G)] * |G|! * (N - |G| - 1)! \quad (1)$$

To use Shapley values in our context, we interpret every DNA read $d \in D$ as a player of a game, and $v(G)$ as the result vector of running Algorithm 1 on a subset $G \subseteq D$. Then, the output of the formula above is a vector that its j 'th entry encodes the read's contribution to the estimated proportion of species j . However, to apply this idea for generating read-level explanations, two main adaptations are needed:

Execution on Samples. Shapley value formula requires executing an algorithm on a subset of its input set. When these values are used in Machine Learning models, usually this requires the output of a model on a subset of its features, to quantify the contribution of the different features to the result of the model. In general Machine Learning models, it is not feasible to calculate the output of the model only on part of the features, as this requires to train the entire model from scratch (a problem that SHAP solves by using the background distribution to sample values for features it integrates out [4]). However, in our Maximum Likelihood estimator, we actually can run the model on a subset of the DNA reads, by simply rewrite the formula in line 15 of Algorithm 1 so that we plug-in a subset G instead of the entire dataset D , and thus calculate what would be the output of the algorithm on the subset G . Furthermore, when we consider multiple subsets G , we only need to execute the expensive PreProcessing step once, and re-use the $A_{s,d}$ values it computes in multiple invocations of $\text{MaximizeLikelihood}(A, G)$. This optimization allows our system to evaluate multiple outputs of the algorithm on different subsets of the input dataset, and to use that information to apply Shapley values formula and attribute the output of the model to the different reads in the dataset.

Scaling Shapley values. Since Algorithm 1 outputs proportions, and not absolute values, the Shapley formula needs to be refined, as illustrated in the following example:

Example 3.3. Consider a sample G for which the ground truth proportion is $(\frac{a}{|G|}, \frac{b}{|G|}, \frac{c}{|G|})$. Further consider a read $i \notin G$ generated from a Homo sapiens. The ground truth for $G \cup \{i\}$ is $(\frac{a+1}{|G|+1}, \frac{b}{|G|+1}, \frac{c}{|G|+1})$.

Thus $v(G) - v(G \cup \{i\}) = (\frac{|G|-a}{|G|^2+|G|}, \frac{-b}{|G|^2+|G|}, \frac{-c}{|G|^2+|G|})$ which is equal to $(\frac{b+c}{|G|^2+|G|}, \frac{-b}{|G|^2+|G|}, \frac{-c}{|G|^2+|G|})$.

Note that $\|v(G) - v(G \cup \{i\})\|$ decreases approximately linearly as $|G|$ grows (see Figure 2 for empirical example that shows that the bigger the sample, the smaller the influence of adding a single read on the output).

Thus, before applying Shapley values in our context, there is a need to address the fact that the magnitude of the influence of adding a single data point to a set depends on the size of that set, otherwise the explanations are dominated by the smaller samples.




Figure 2: X axis is number of reads in the sample, Y axis is the average magnitude of change that is caused by adding a single read to a sample. The plot shows empirically that the more reads are in a sample, the smaller the influence of adding a single read is.

We therefore added a scaling parameter to the original Shapley value formula, that takes into account the size of the subset. Specifically, to generate a read-level explanation for read i , exML samples M subsets of D , and outputs the following vector as explanation (This can be calculated efficiently, as described above, since the preprocessing step only needs to be executed once):

$$\frac{1}{M} \sum_{G \in \text{Samples}} [v(G \cup i) - v(G)] * |\mathbf{G}|. \quad (2)$$

Given a read $d \in D$, the output of this process is an explanation vector v^d that corresponds to definition 3.1, and this is the vector that our system generates as an attribution based read-level explanation.

Example 3.4 (Scaled Shapley values). Figure 3 shows read-level explanations generated by exML using the scaled Shapley values on D_1 and D_2 (as described in example 2.1). As stated above, the output of Algorithm 1 on both datasets is approximately the same, so using merely the output of Algorithm 1, one may not distinguish between them. However, as Figure 3 demonstrates, their different read-level explanations provide insights that direct the user to the correct interpretation of the results. These explanations reflect that in D_1 , only the last 40 reads had substantial influence on the output, whereas in D_2 , all reads were influential. This implies that the system has more evidence to its output on D_2 . In addition, the non-influential reads in D_1 might also be attributed to being originated from a genomically distant organism that is not in the reference set and hence do not influence the output proportion of neither Homo sapiens nor Denisovans.

3.1.2 Read-level counterfactuals. Given a model M and an input x such that $M(x) = l$, a *counterfactual explanation* is a modified instance y for which $M(y) \neq l$. Intuitively, modifications to x that lead to a different label are indicative of the features that were responsible for the label l being chosen to begin with. We now define and elaborate on the 3 types of counterfactual our system produces. We start by defining few useful notations in the settings of Maximum Likelihood, and then use these notations to define the counterfactual explanations the system produces.

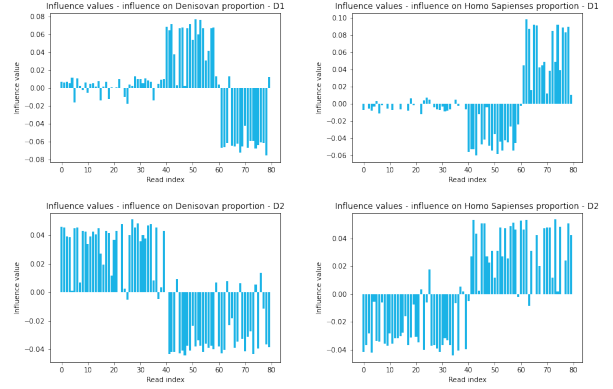


Figure 3: Read-level explanations on D_1 (top), and D_2 (bottom). x axis is the read index, and y axis is the explanation value. Right plots show the explanation values of Homo Sapiens proportions, and Left plots show explanation values on Denisovan proportions.

Definition 3.5. Given a dataset D , a set of labels S and a likelihood function $L(\alpha, D) \in ([0, 1]^{|S|}, \text{set}) \mapsto [0, 1]$, define $ML(D, L) := \text{argmax}_{\alpha} L(\alpha, D)$.

Intuitively, the likelihood function gets as input a distribution vector over the labels and a set, and outputs a probability. For example, in our settings, if D is a set of DNA reads, α is a proportion vector, S is the set of known hominin species, and L is the likelihood function of Algorithm 1 on the dataset D , $ML(D, L)$ is the vector that maximizes the likelihood term $L(\alpha, D)$, for example, it can be the vector $(0, 0, 1)$. This is typically the output of the maximum likelihood algorithm.

Definition 3.6. Let $H(D, L)$ be $\text{argmax}_i (ML(D, L)[i])$ - the index of the maximal value in the vector $ML(D, L)$.

Intuitively $H(D, L)$ is the label with the highest proportion assigned to it. In our example, if $ML(D, L) = (0, 0, 1)$, so $H(D, L) = 2$ as this is the index that maximizes $ML(D, L)$. In our example this corresponds to saying that the algorithm estimates that the Denisovans species is the most abundant in D .

Definition 3.7. Let *changing dominance subset* be a subset $D' \subseteq D$ such that $H(D, L) \neq H(D \setminus D', L)$.

This is a subset that removing it from the input dataset causes the vector that maximizes L to have a different maximizing index. In our example, if there are 2 DNA reads $D_1, D_2 \in D$ that removing them from D will cause the output of Algorithm 1 to be maximal when $\alpha = (1, 0, 0)$, then the set $\{D_1, D_2\}$ is a changing dominance subset with respect to L and D .

Definition 3.8. Define $CF1(L, D)$ as a subset $D' \subseteq D$ that is a changing dominant subset, and is also minimum subset with that property, meaning that for every subset $T \subseteq D$ that is changing dominance with respect to L , $|D'| \leq |T|$.

In DNA classification, this counterfactual is an interesting research insight, as it points out subsets of the data that make the

algorithm substantially change its decision, which provides information both on these subsets (maybe they originate from different species then the others) and on the stability of the output (if a relatively small subset is sufficient to make a substantial change to the output, the algorithm is less stable and trustworthy).

If we assume that L can be any function, even if we have an $O(1)$ oracle for the dominant species given a subset, the problem is NP-hard.

Unfortunately, the exact calculation of CF1 given a function and a dataset is NP-hard. We will now describe a reduction that shows that this problem is NP-hard, and then present a greedy algorithm that we created and that our system uses to output estimation to CF1.

reduction from SAT. Given a boolean formula ψ in variables $x_1 \dots x_n$, we will build an instance of the counterfactual problem and show that finding the counterfactual subset is equivalent to finding a satisfying assignment to ψ .

Let D be the set $\{1 \dots n + 1\}$.

$L(\alpha, D)$ will be 1 if $\alpha = (1, 0, 0)$ and 0 otherwise, so the dominant species on all the dataset is species 0.

$L(\alpha, D')$ will be equal to $L(\alpha, D)$ if D' corresponds to not satisfying assignment, and will be the opposite if it is satisfying. An assignment is converted to a group simply by taking into the group the variables that are getting true in the assignment. The reduction is polynomial time since it is just checking the assignment (given a formula it builds the function L that only needs to convert the assignment to a subset of D and to check the assignment on the formula ψ). If someone solves the counterfactual problem and finds a subset that is changing dominance, it means that we can take the changing dominance subset and convert it to a satisfying assignment, and solve SAT. If there is no satisfying assignment, then there will be no changing dominance subset.

algorithm 2 is a greedy algorithm that aims to estimate cf1. It is based on removing reads from the dataset that have the Shapley value that is most in favor of the dominant species, hoping to quickly converge to a dataset in which this is no longer the dominant species.

Algorithm 2: Calculating counterfactual 1

Data: dataset D , and the output of Algorithm 1 on that dataset

Result: subset that is an estimation of cf1

- 1 Assume w.l.o.g that the dominant species in D is homo sapiens
 - 2 Calculate scaled shapley values for every $d \in D$, define v^d as the vector of scaled shapely value of read d , and v_{hs}^d is the influence of this read on homo sapiens proportion.
 - 3 Start with $D' = \emptyset$, and add to it the read d' , which is $\text{argmax}_{d \notin D'} (v_{hs}^d)$
 - 4 Repeat until $H(D \setminus D', L) \neq \text{Homosapiens}$
 - 5 **return** D' as an estimation to the minimal changing dominance subset
-

Change dominating species from 'Homo Sapiens' to 'Neanderthal' would require removing 11 reads:
[42, 57, 54, 52, 58, 56, 43, 46, 45, 44, 59]

Change dominating species from 'Homo Sapiens' to 'Neanderthal' would require removing 30 reads:
[32, 34, 38, 49, 37, 57, 22, 46, 51, 42, 50, 58, 23, 55, 53, 59, 41, 36, 25, 39, 26, 44, 40, 29, 15, 2, 47, 56, 5, 24]

Figure 4: CF1 output on D_3 (top) and D_4 (bottom)

In every iteration, the algorithm is removing from the dataset the read that is mostly "supporting" the dominant species proportion, expecting to quickly obtain a dataset in which the dominant species is different. It is easy to see that the set that is returned is changing dominance, and that if there exists a changing dominance subset, this algorithm will return a changing dominance subset. But as stated above, this is only an approximation to the *minimum* changing dominance subset.

CF2 is a minimum subset $D' \subset D$ that changes the model's output by at least ϵ , i.e. $\|ML(D) - ML(D \setminus D')\| > \epsilon$. To estimate it, we initialize $D' = \emptyset$, and greedily add to D' a read d that maximizes the term $\| \sum_{i \in D'} S(i) + S(d) \|$, until $\|v(D) - v(D \setminus D')\| > \epsilon$. Note that since we are comparing vector norms, d is not necessarily the read with the highest norm of the Shapley vector, but rather one that shifts the decision in the right direction w.r.t. D' .

CF3 is a maximum subset $D'_s \subseteq D$ such that $v(D'_s)[s] \geq \theta$ for each $s \in S$, where $v(D)[s]$ is the estimated proportion that Algorithm 1 assigns to species s when run on D . To estimate CF3 for a species s , we initialize $D'_s = \emptyset$, and greedily add a read d that has the highest $S(d)$ w.r.t. s , until there are no reads for which $v(D'_s \cup d)[s] \geq \theta$.

Example 3.9 (Read-level counterfactuals). Figure 4 shows the output of CF1 on two datasets: D_3 (top figure) has 20 reads from Homo sapiens, and 40 non-informative reads, and D_4 (bottom) has 40 reads from Homo Sapiens and 20 non-informative reads. Though Algorithm 1's output is the same for both datasets $((1, 0, 0))$, Figure 4 shows that on D_3 , CF1 contains 11 reads, whereas in D_4 , CF1 contains 30 reads. This conveys to the user that the output on D_4 is more stable, as a larger modification is required for it to change.

3.2 Reference-level explanations

We are interested in generating explanations for references - to quantify and explain to the user how every reference in the labeled references set influenced the output of Algorithm 1. We present Algorithm 3, a sampling based algorithm to calculate explanation vectors for references. The explanations that we compute follow the attribution approach. We define $A_{s,d}^i$ as the value that Algorithm 1 would calculate as $A_{s,d}$, if the references set would not include r_i . In line 1, we run the PreProcessing method of Algorithm 1, to get the initial $A_{s,d}$ values. In line 2 we loop over the DNA references. For $A_{s,d}^i$ values where s is the label of r_i , we set the $A_{s,d}^i$ value to be the average that would have been obtained if r_i was not in the references set (line 3); otherwise, we set $A_{s,d}^i = A_{s,d}$ (line 4). In lines 5–9 we go over the samples, and calculate the change that is caused to the output vector of Algorithm 1 by removing r_i . In line 10 we average this influence value over all samples, to get an estimation of how removing reference i influences the output of Algorithm 1. In

line 12, we return V , which is a list of all V_i values (V_i is the vector exML outputs as an explanation to reference i , and it intuitively describes how reference i influences the output of Algorithm 1).

Algorithm 3: Calculating reference-level explanations

Data: References set R , input dataset D , Samples K , such that every $K_j \in K$ is a subset of D

Result: V , s.t V_i is the reference-level explanation of R_i

```

1 Run preProcessing() step of algorithm 1 to get  $A_{s,d}$  and  $p[i,d]$  values
2 foreach  $r_i \in R$  do
3   For each  $A_{s,d}$  such that  $s = Y_i$ ,  $A_{s,d}^i := \frac{A_{s,d} * |R_s| - p[r_i,d]}{|R_s| - 1}$ 
4   For each  $A_{s,d}$  such that  $s \neq Y_i$ ,  $A_{s,d}^i := A_{s,d}$ 
5   foreach  $k_j \in K$  do
6      $L^i(k_j; \alpha) := \prod_{d \in k_j} \sum_{s \in S} \alpha_s * A_{s,d}^i$ 
7      $L(k_j; \alpha) := \prod_{d \in k_j} \sum_{s \in S} \alpha_s * A_{s,d}$ 
8      $\delta_{i,j} \leftarrow \text{argmax}_{\alpha} (L^i(k_j; \alpha)) - \text{argmax}_{\alpha} (L(k_j; \alpha))$ 
9   end
10   $V_i = \frac{\sum_{j \in [K]} \delta_{i,j}}{|K|}$ 
11 end
12 return  $V$ 

```

Example 3.10. Figure 5 shows the reference explanations generated by our system on two datasets, D_5 (left figure), and D_6 (right figure). Each dataset contains 50 reads generated from Neanderthals, 50 from Homo sapiens, and 50 from Denisovans. However, D_5 's Neanderthals are from Neanderthal KX198082, and D_6 's Neanderthals are from Neanderthal Altai KC879692. On both datasets, Algorithm 1 outputs the same estimation ($\frac{1}{3}, \frac{1}{3}, \frac{1}{3}$), but the reference explanations can help identify the specific sub group of Neanderthals that contributed to the dataset, as in both explanations, the reference that has the highest positive influence on the Neanderthal proportion is the one from which the data was actually generated.

4 IMPLEMENTATION

We implemented *exML* using python, while utilizing the following tools: *biopython* for alignments and parsing of DNA data, *joblib* for parallel computation, and *SHAP* for generating visualizations. Reference genomes were downloaded from the *NCBI database*, and simulated ancient DNA data was generated using a python script published in [7].

The source code can be found here, together with instructions on how to run the code, or to load and run a GUI we created that interacts with the system.

REFERENCES

- [1] Ewan Birney, Jessica Vamathevan, and Peter Goodhand. 2017. Genomics in healthcare: GA4GH looks to 2022. *BioRxiv* (2017), 203554.
- [2] Solomon Kullback and Richard A Leibler. 1951. On information and sufficiency. *The annals of mathematical statistics* 22, 1 (1951), 79–86.
- [3] Yichen Liu, Xiaowei Mao, Johannes Krause, and Qiaomei Fu. 2021. Insights into human history from the first decade of ancient genomics. *Science* 373, 6562 (2021), 1479–1484.
- [4] Scott M Lundberg and Su-In Lee. 2017. A Unified Approach to Interpreting Model Predictions. In *Advances in Neural Information Processing Systems*, I. Guyon, U. V.

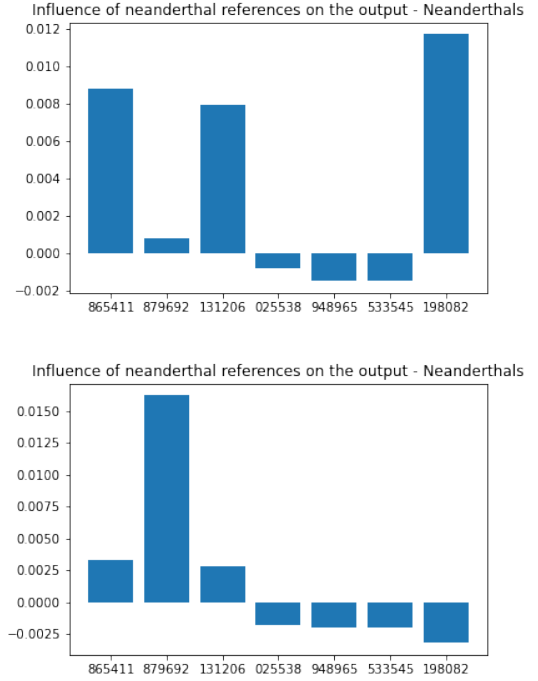


Figure 5: Reference explanations on D_5 (top) and D_6 (bottom). X axis is the reference index, and y axis is the influence of that reference on the proportion of Neanderthals in the output of algorithm 1.

- Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett (Eds.), Vol. 30. Curran Associates, Inc.
- [5] Alice Carolyn McHardy, Héctor García Martín, Aristotelis Tsirigios, Philip Hugenoltz, and Isidore Rigoutsos. 2007. Accurate phylogenetic classification of variable-length DNA fragments. *Nature methods* 4, 1 (2007), 63–72.
- [6] J Craig Venter, Mark D Adams, Eugene W Myers, Peter W Li, Richard J Mural, Granger G Sutton, Hamilton O Smith, Mark Yandell, Cheryl A Evans, Robert A Holt, et al. 2001. The sequence of the human genome. *science* 291, 5507 (2001), 1304–1351.
- [7] Benjamin Vernot, Elena I Zavala, Asier Gómez-Olivencia, Zenobia Jacobs, Viviane Slon, Fabrizio Mafessoni, Frédéric Romagné, Alice Pearson, Martin Petr, Nohemi Sala, et al. 2021. Unearthing Neanderthal population history using nuclear and mitochondrial DNA from cave sediments. *Science* 372, 6542 (2021), eabf1667.
- [8] David S Watson. 2021. Interpretable machine learning for genomics. *Human genetics* (2021), 1–15.